



# MULTICOLUMN INDICES, JOINS & AGGREGATIONS

---

# MULTICOLUMN INDICES

# MULTICOLUMN INDEX (COMPOSITE INDEX)

---

```
create index index_documents_on_stp on  
documents(supplier, total_amount,  
published_on);
```

{ ANAsoft | 1000 | 12. 9. 2019 }

{ ANAsoft | 1000 | 12. 10. 2019 }

{ ANAsoft | 2000 | 1. 1. 2018 }

{ BETHESSDA | 300 | 18. 11. 2009 }

{ BETHESSDA | 1000 | 16. 5. 2010 }

{ BETHESSDA | 1000 | 17. 5. 2010 }

.. ..



```
explain select * from documents where supplier =  
'SPP';
```

**Bitmap Heap Scan** on documents (cost=4.48..35.99  
rows=8 width=303)

  Recheck Cond: ((supplier)::text = 'SPP'::text)

  -> **Bitmap Index Scan** on  
index\_documents\_on\_stp (cost=0.00..4.48 rows=8  
width=0)

    Index Cond: ((supplier)::text =  
'SPP'::text)

```
explain select * from documents where supplier =  
'SPP' and total_amount = 2000;
```

**Index Scan** using index\_documents\_on\_stp on  
documents (cost=0.42..8.44 rows=1 width=303)  
  Index Cond: (((supplier)::text = 'SPP)::text)  
AND (total\_amount = 2000::double precision))

```
explain select * from documents where supplier =  
'SPP' and total_amount = 2000 and published_on =  
now();
```

**Index Scan** using index\_documents\_on\_stp on  
documents (cost=0.42..8.45 rows=1 width=303)  
  Index Cond: (((supplier)::text = 'SPP)::text)  
AND (total\_amount = 2000::double precision) AND  
(published\_on = now()))

# NON-LEADING COLUMN CONDITIONS

---

```
explain select * from documents where  
total_amount = 2000;
```

**Seq Scan** on documents (cost=0.00..17398.30  
rows=1276 width=303)  
Filter: (total\_amount = 2000::double  
precision)

# NON-LEADING COLUMN CONDITIONS - INDEX SCAN ?!

---

```
explain select * from documents where total_amount  
= 2000 and published_on > '2012-01-01';
```

**Bitmap Heap Scan** on documents

(cost=15348.86..17840.69 rows=779 width=303)

  Recheck Cond: ((total\_amount = 2000::double precision) AND (published\_on > '2012-01-01 00:00:00'::timestamp without time zone))

  -> Bitmap Index Scan on index\_documents\_on\_stp  
(cost=0.00..15348.66 rows=779 width=0)

    Index Cond: ((total\_amount = 2000::double precision) AND (published\_on > '2012-01-01 00:00:00'::timestamp without time zone))

# BENCHMARK

---

```
explain analyze select * from documents where  
total_amount = 2000 and published_on > '2012-01-01';
```

Bitmap Heap Scan on documents  
(cost=15348.86..17840.69 rows=779 width=303) (actual  
time=19.163..19.663 rows=840 loops=1)  
  Recheck Cond: ((total\_amount = 2000::double  
precision) AND (published\_on > '2012-01-01  
00:00:00'::timestamp without time zone))  
    -> Bitmap Index Scan on index\_documents\_on\_stp  
(cost=0.00..15348.66 rows=779 width=0) (actual  
time=19.076..19.076 rows=840 loops=1)  
      Index Cond: ((total\_amount = 2000::double  
precision) AND (published\_on > '2012-01-01  
00:00:00'::timestamp without time zone))  
**Total runtime: 19.739 ms**

# BENCHMARK

---

```
set enable_indexscan = off;
set enable_bitmapscan = off;
explain analyze select * from documents where
total_amount = 2000 and published_on >
'2012-01-01';
```

Seq Scan on documents (cost=0.00..18276.36  
rows=779 width=303) (actual time=29.384..78.961  
rows=840 loops=1)

  Filter: ((published\_on > '2012-01-01  
00:00:00'::timestamp without time zone) AND  
(total\_amount = 2000::double precision))

  Rows Removed by Filter: 350384

Total runtime: **79.021 ms**

## 3 SEPARATE SINGLE-COLUMN INDICES

---

```
drop index index_documents_on_stp;
```

```
create index index_documents_on_supplier on  
documents(supplier);
```

```
create index index_documents_on_total_amount on  
documents(total_amount);
```

```
create index index_documents_on_published_on on  
documents(published_on);
```

```
.....  
explain select * from documents where supplier = 'SPP'  
and total_amount = 2000;
```

**Bitmap Heap Scan** on documents (cost=30.73..34.74 rows=1  
width=303)

  Recheck Cond: (((supplier)::text = 'SPP'::text) AND  
(total\_amount = 2000::double precision))

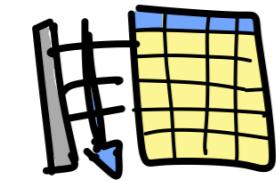
  -> **BitmapAnd** (cost=30.73..30.73 rows=1 width=0)

    -> **Bitmap Index Scan** on  
index\_documents\_on\_supplier (cost=0.00..4.48 rows=8  
width=0)

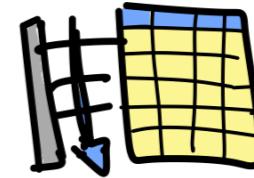
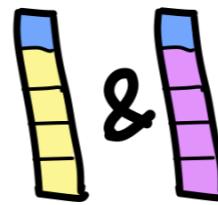
      Index Cond: ((supplier)::text =  
'SPP'::text)

      -> **Bitmap Index Scan** on  
index\_documents\_on\_total\_amount (cost=0.00..25.99  
rows=1276 width=0)

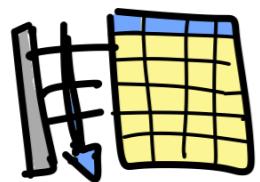
      Index Cond: (total\_amount = 2000::double  
precision)



SUPPLIER



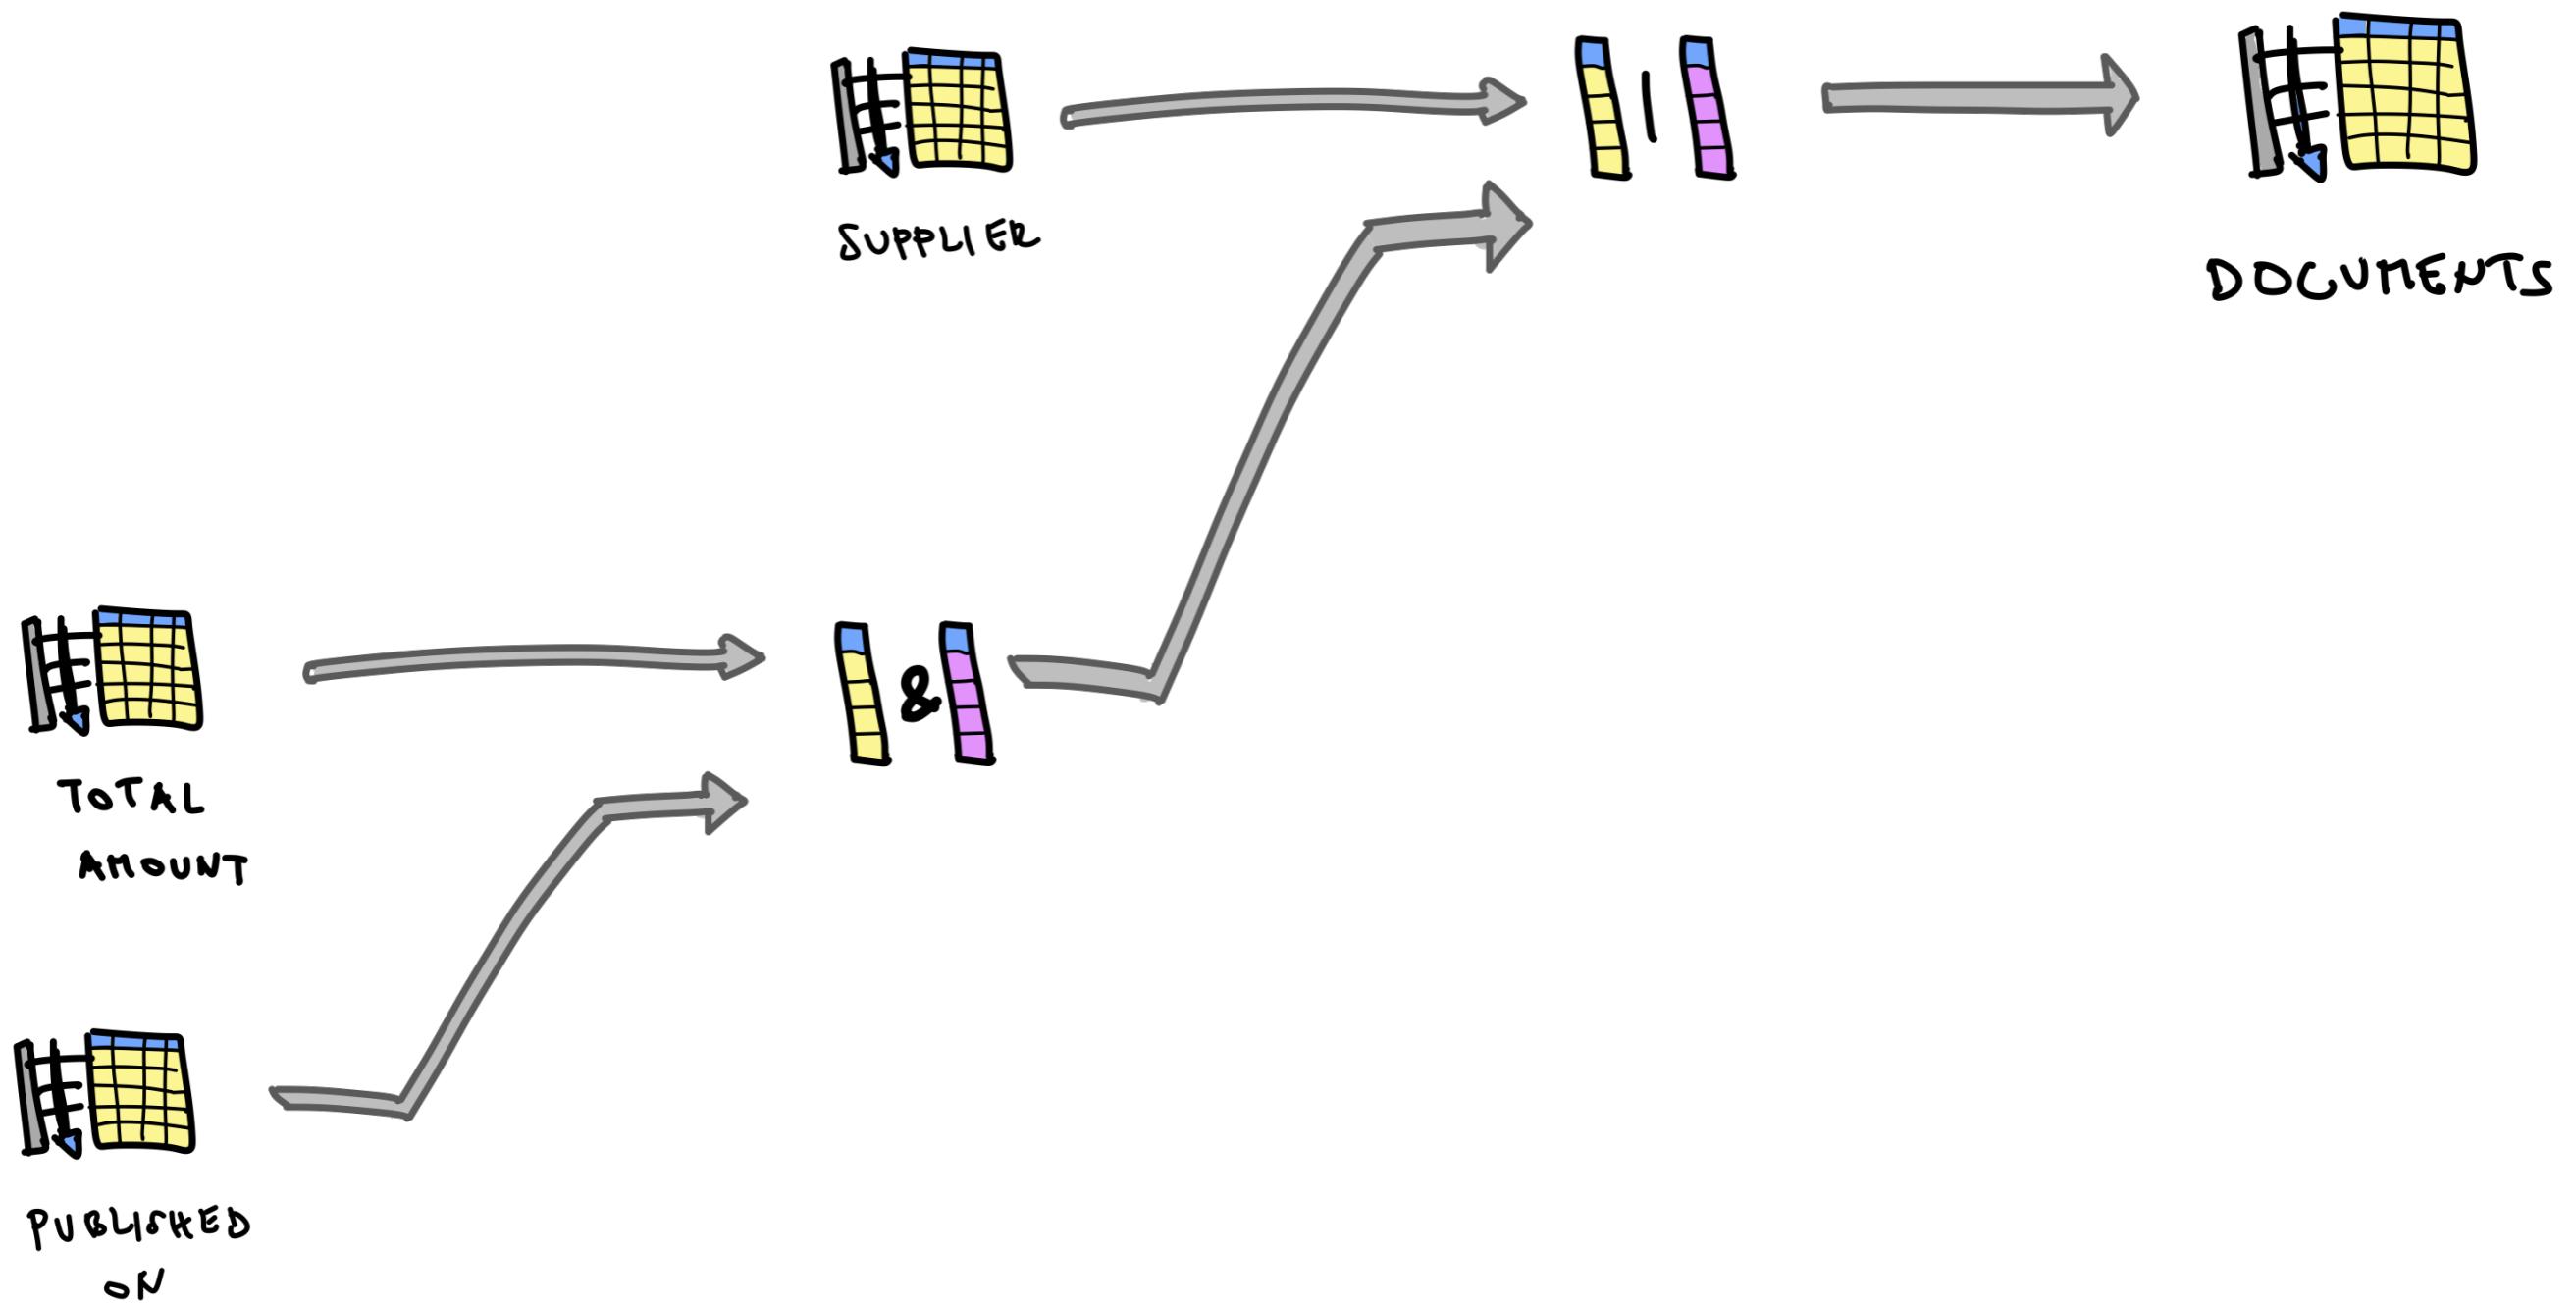
DOCUMENTS



TOTAL  
AMOUNT



```
.....  
select *  
  from documents  
where supplier = 'SPP'  
  or total_amount = 2000  
and published_on > '2013-01-01';
```



# COMPOSITE INDEX VS. PER-COLUMN INDEX

---

- Per-column indices:
  - more flexible
  - support any column combinations in the *where* clause
- Use composite index for:
  - efficient *order by* on 2 and more columns
  - covering index - piggybacking data (add data to index to leverage index-only scan and avoid heap scans)

# COVERING INDEX

---

```
create index index_documents_on_dc on
documents(department, customer);

vacuum analyze;

explain select customer from documents where
department = 'Rozhlas a televizia Slovenska';

Index Only Scan using index_documents_on_dc on
documents (cost=0.55..1200.50 rows=16683
width=41)

    Index Cond: (department = 'Rozhlas a televizia
Slovenska'::text)
```

# JOINS

# JOINS

---

- Nested loop
- Hash join
- Merge join

# NESTED LOOP

---

- For each row in table A
  - For each row in table B
    - If join condition is true, emit row (e.g., A.id = B.a\_id)
- Most generic join algorithm, has no expectations

# NESTED LOOP WITH INDEX

---

```
set enable_hashjoin=off;
```

```
explain select * from character c  
  join paragraph p on p.charid = c.charid;
```

**Nested Loop** (cost=0.28..13080.14 rows=35465 width=416)

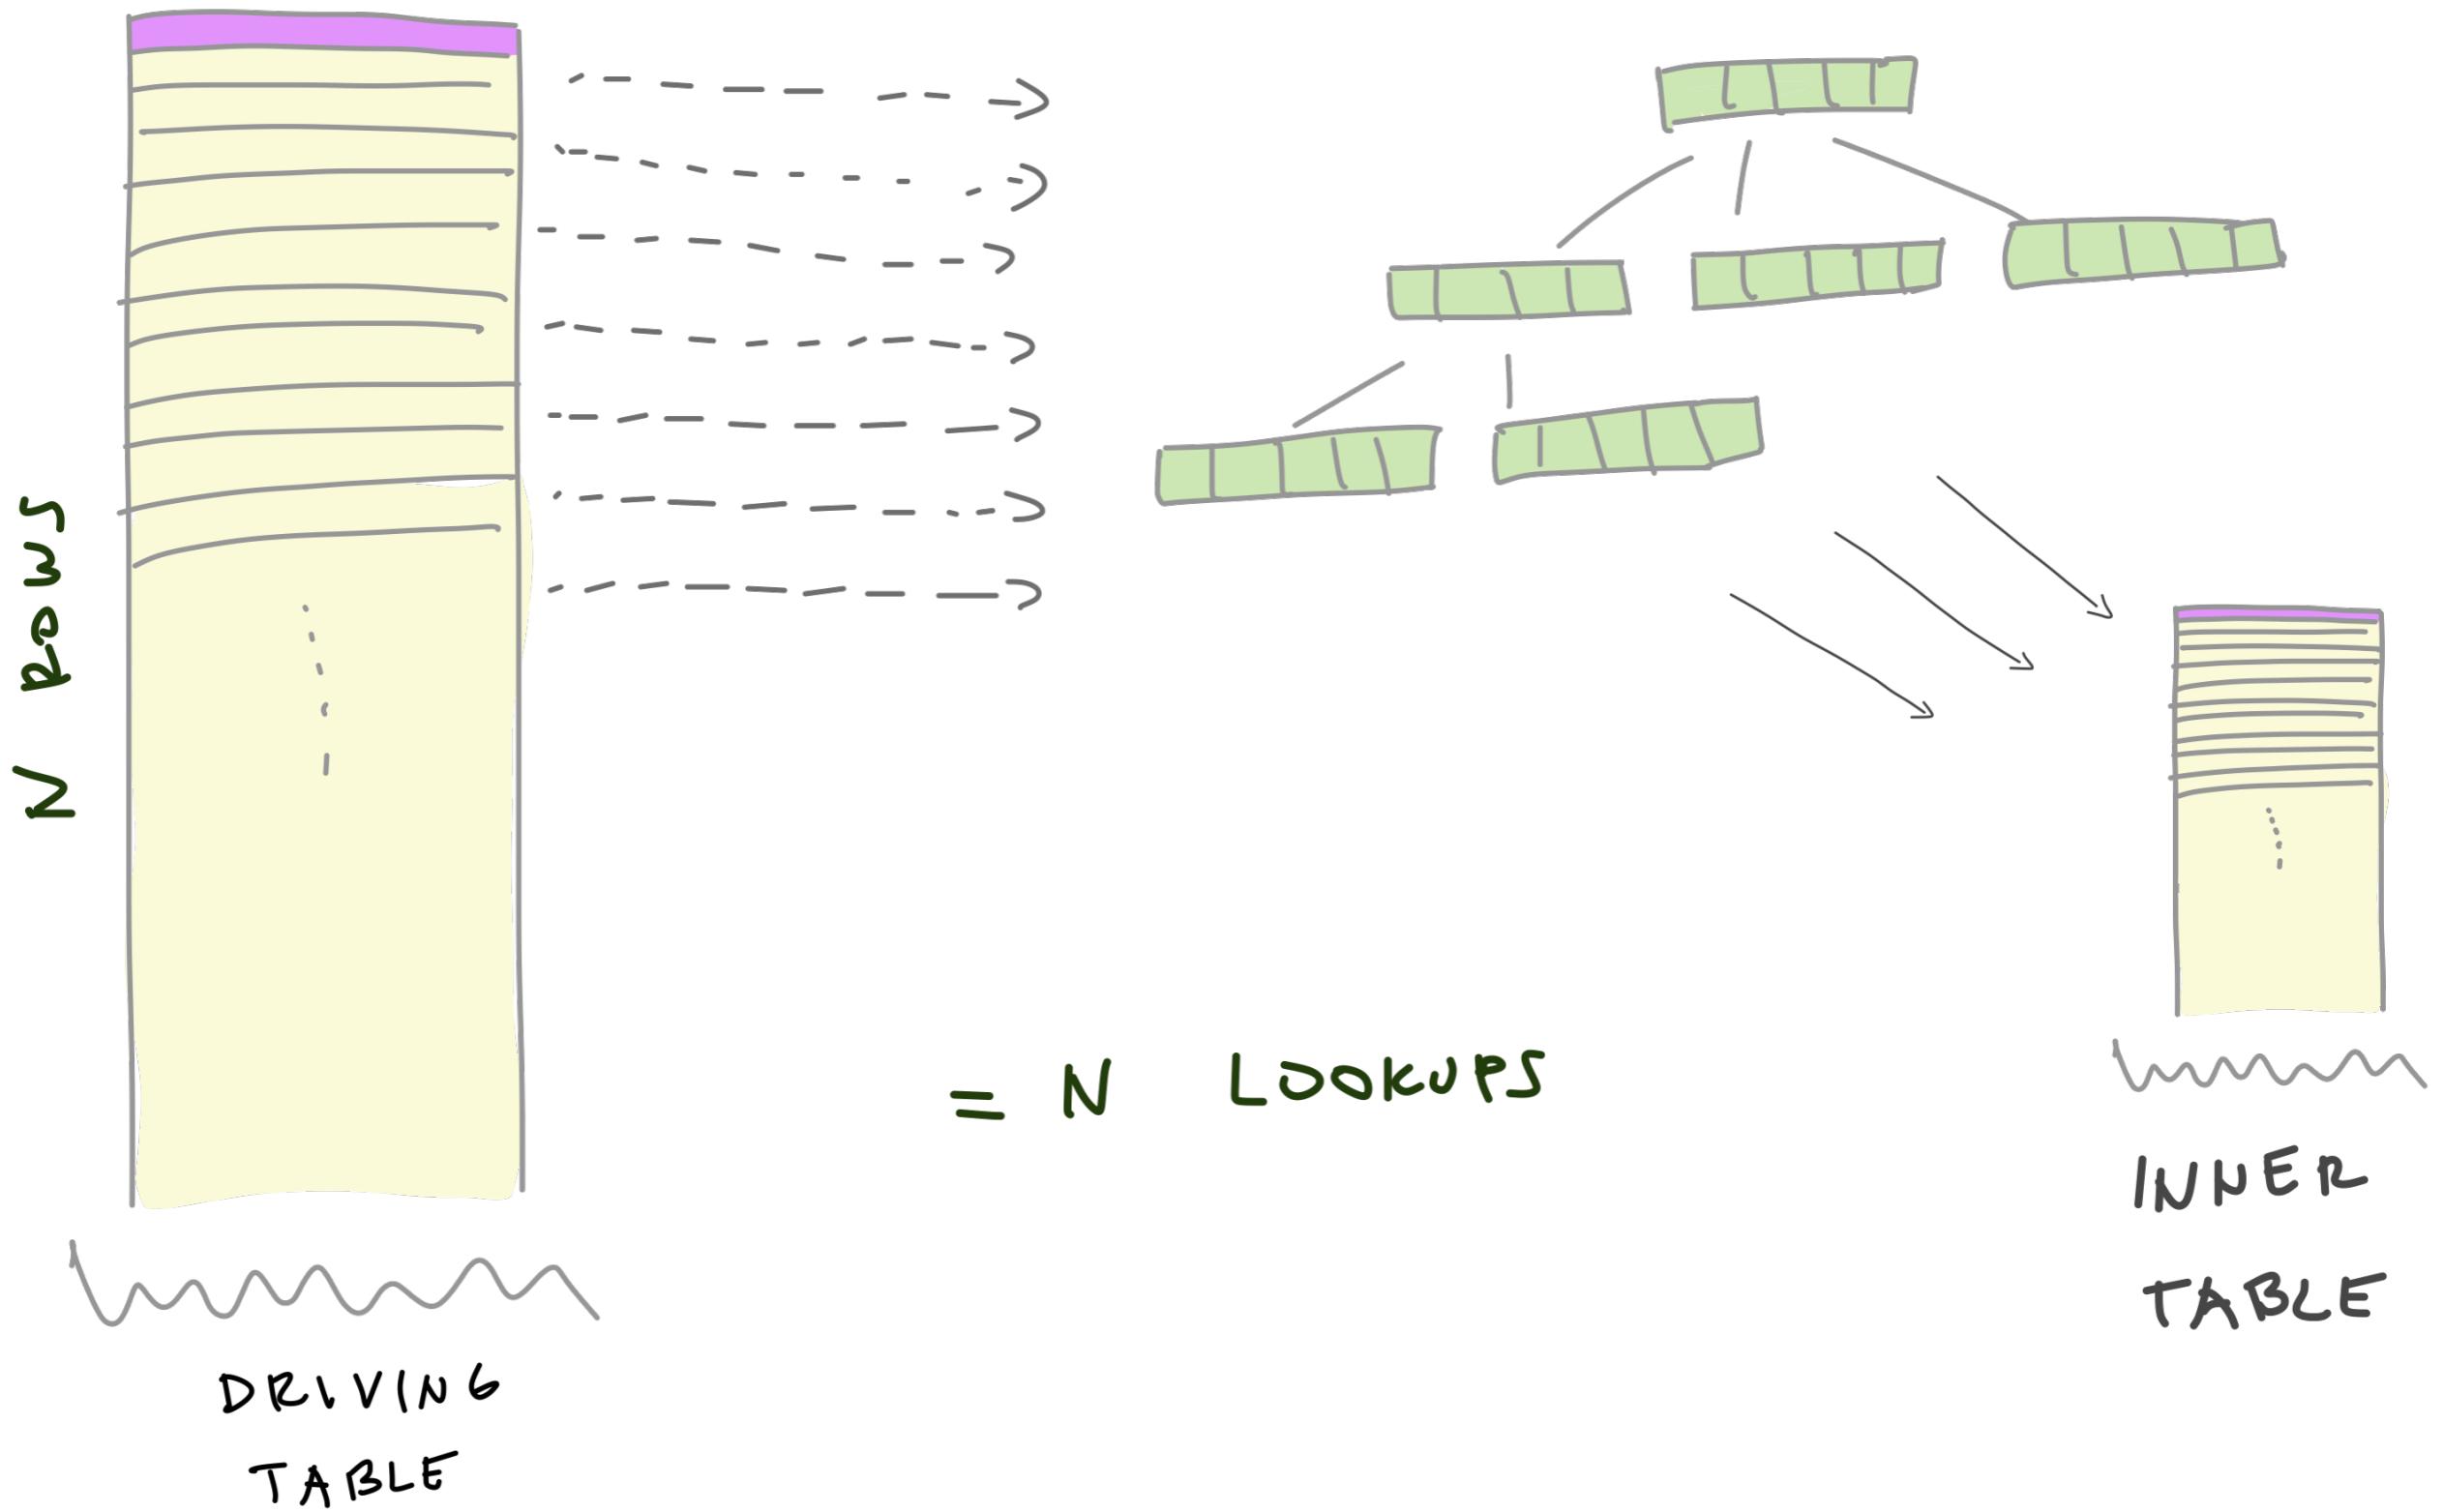
  -> **Seq Scan** on paragraph p  
(cost=0.00..2094.65 rows=35465 width=361)

  -> **Index Scan** using character\_pkey on  
"character" c (cost=0.28..0.30 rows=1 width=55)

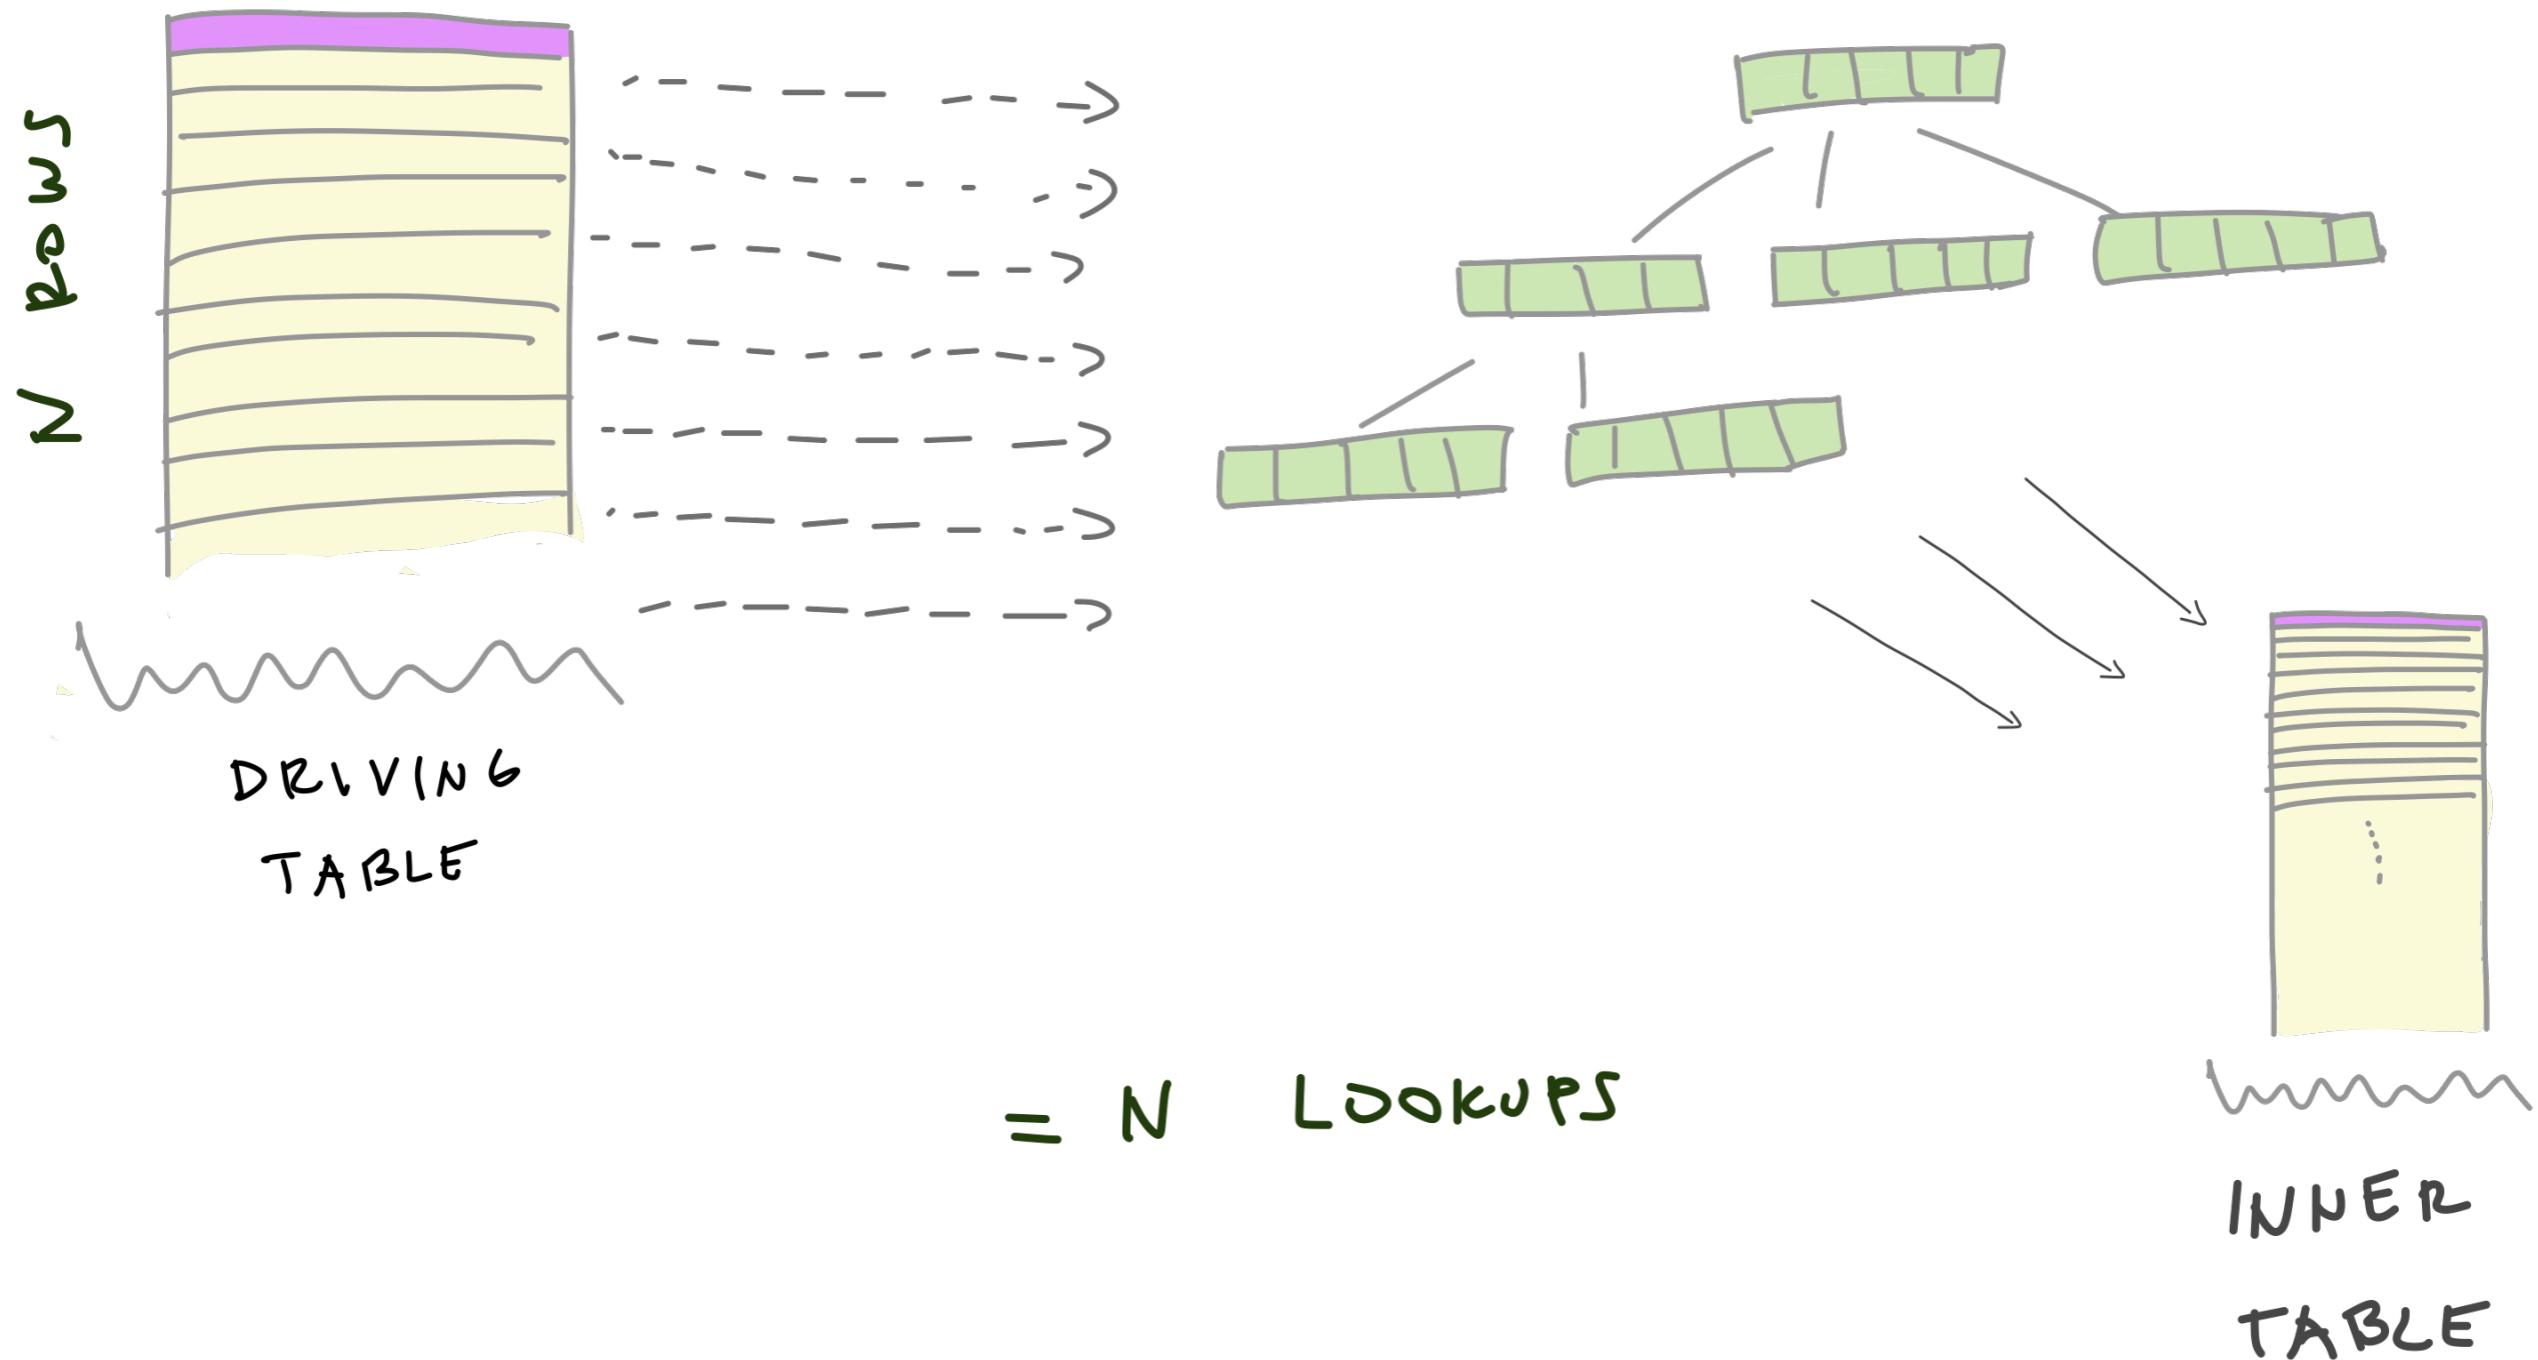
    Index Cond: ((charid)::text =  
(p.charid)::text)

# NESTED LOOPS WITH INDEX

---



# NESTED LOOPS WITH INDEX



# NESTED LOOPS – SUMMARY

---

- The smaller the driving table, the less lookups
- Query planner will make the correct decision based on the statistics. Ignore advice that you need to handle this yourself. (You do, if your database is using RBO, though).
- Nested loops are feasible with index
  - Seq scan for each row of the inner table is.. inefficient
    - CBO will try to avoid it

# HASH JOIN

---

```
set enable_hashjoin=on;
```

```
explain select * from character c  
join paragraph p on p.charid = c.charid;
```

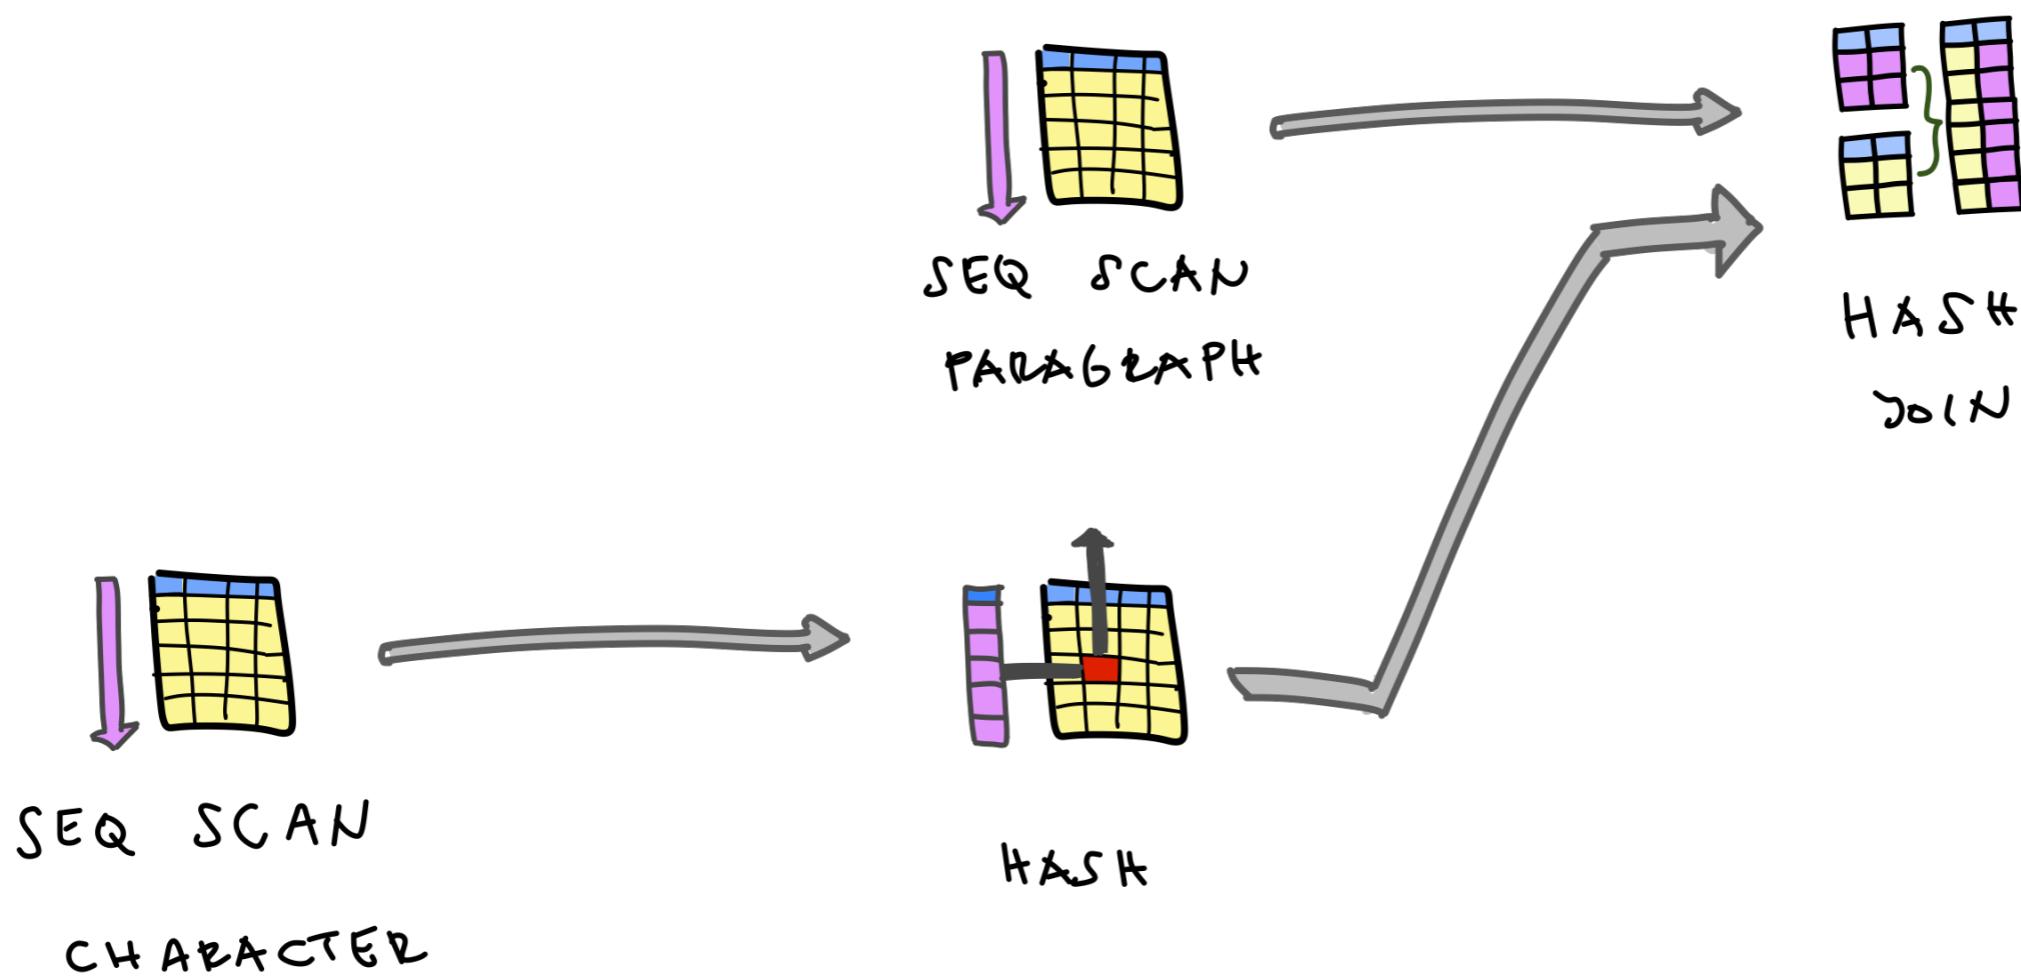
```
Hash Join (cost=41.48..2623.78 rows=35465  
width=416)
```

```
  Hash Cond: ((p.charid)::text = (c.charid)::text)
```

```
    -> Seq Scan on paragraph p (cost=0.00..2094.65  
rows=35465 width=361)
```

```
    -> Hash (cost=25.66..25.66 rows=1266 width=55)
```

```
          -> Seq Scan on "character" c  
(cost=0.00..25.66 rows=1266 width=55)
```



# HASH JOIN

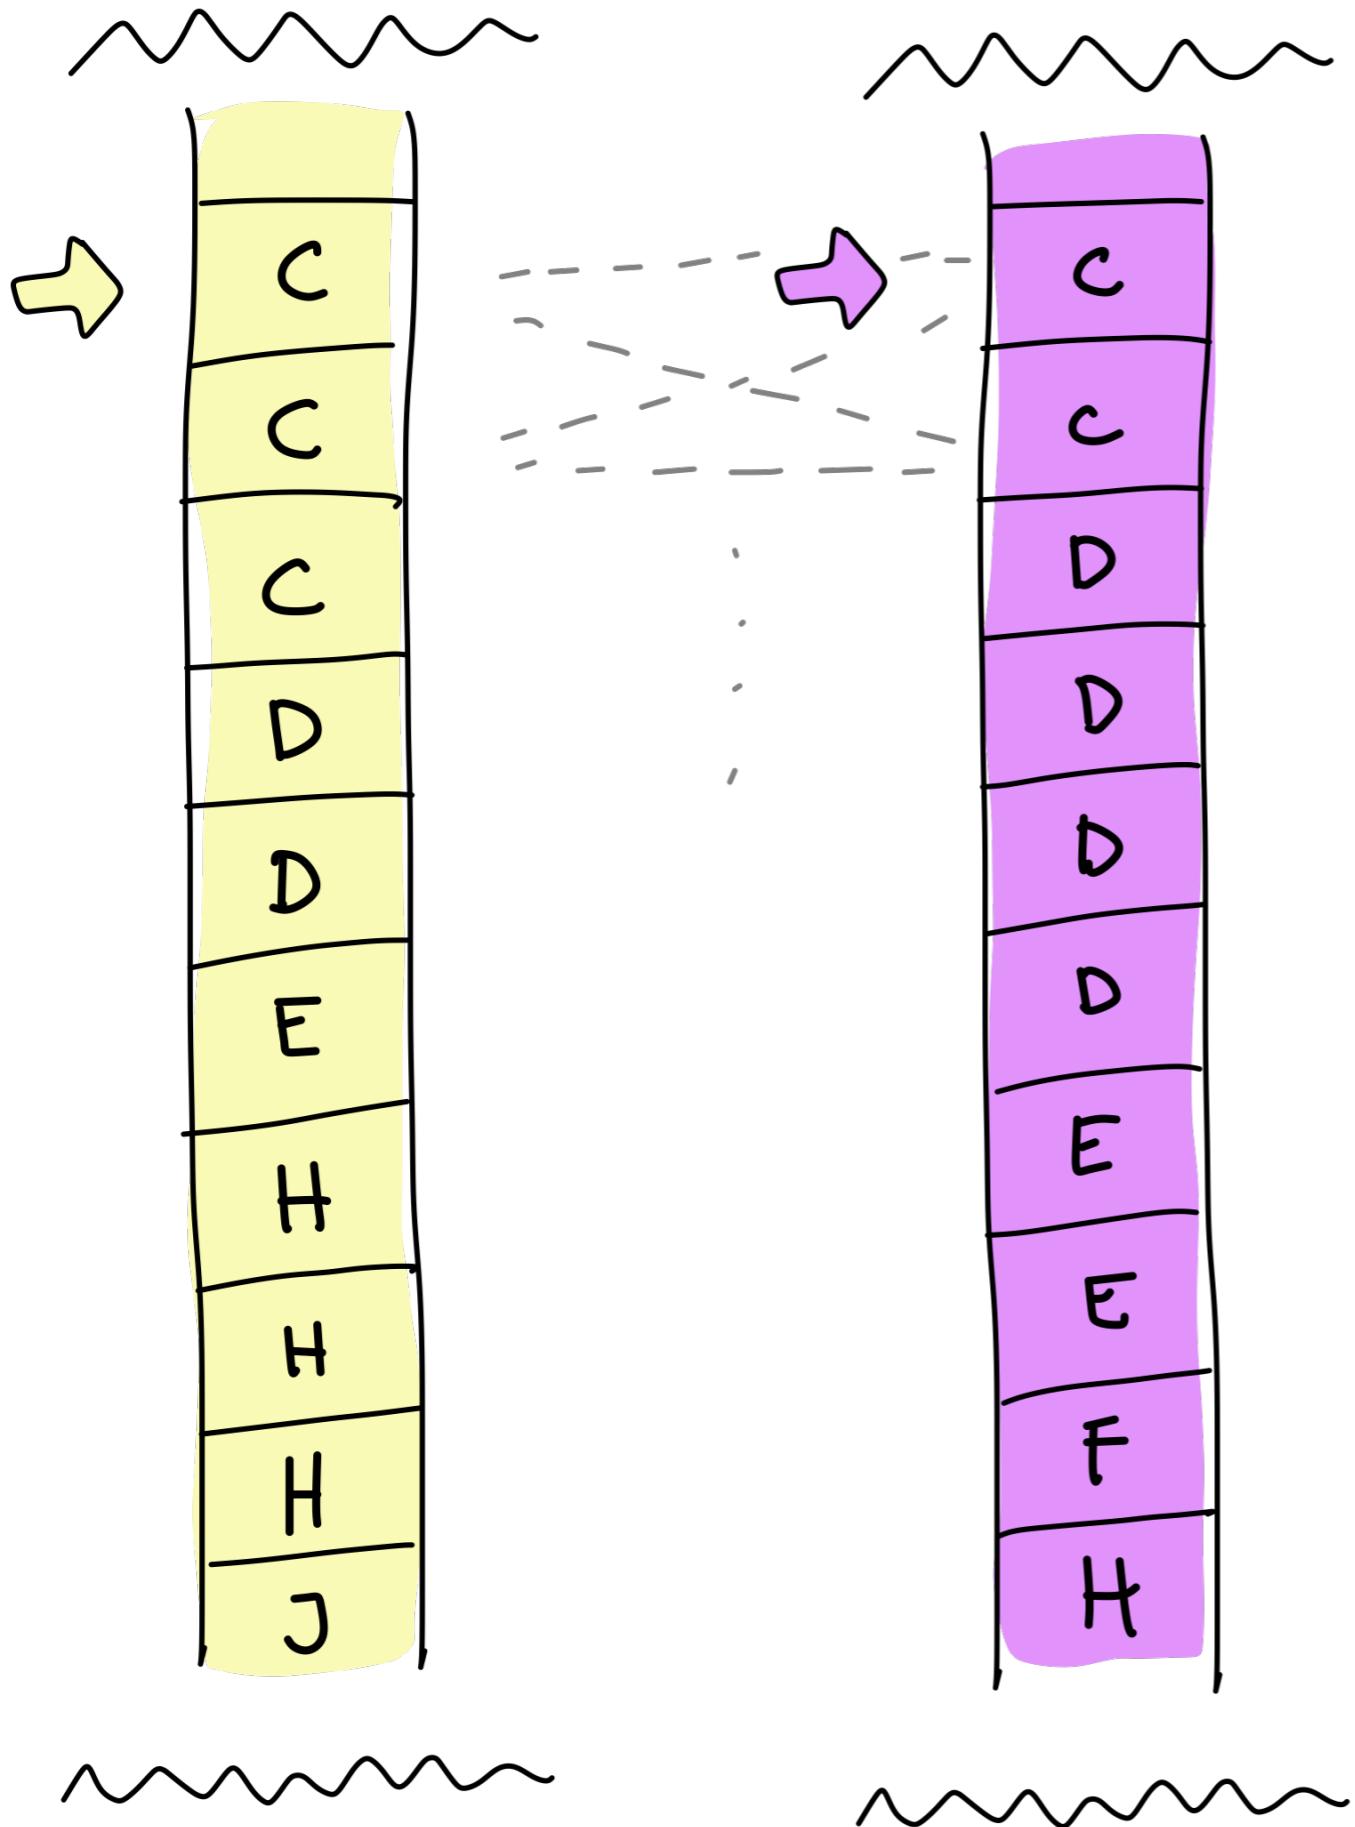
---

- Compute hash table from inner table (hash key = join column)
- For each row in driving table, lookup value in the hash table (lookup key = the other join column)
- Only works for simple equality joins

# MERGE JOIN

---

- Zip algorithm
- Requires both tables to be sorted (only viable via index scan)



# MERGE JOIN

---

```
create index idx_paragraph_on_charid on
paragraph(charid);
```

```
explain select * from character c
  join paragraph p on p.charid = c.charid;
```

```
Merge Join  (cost=0.57..8423.73 rows=35465
width=416)
```

```
  Merge Cond: ((c.charid)::text =
(p.charid)::text)
```

```
    -> Index Scan using character_pkey on
"character" c  (cost=0.28..63.36 rows=1266
width=55)
```

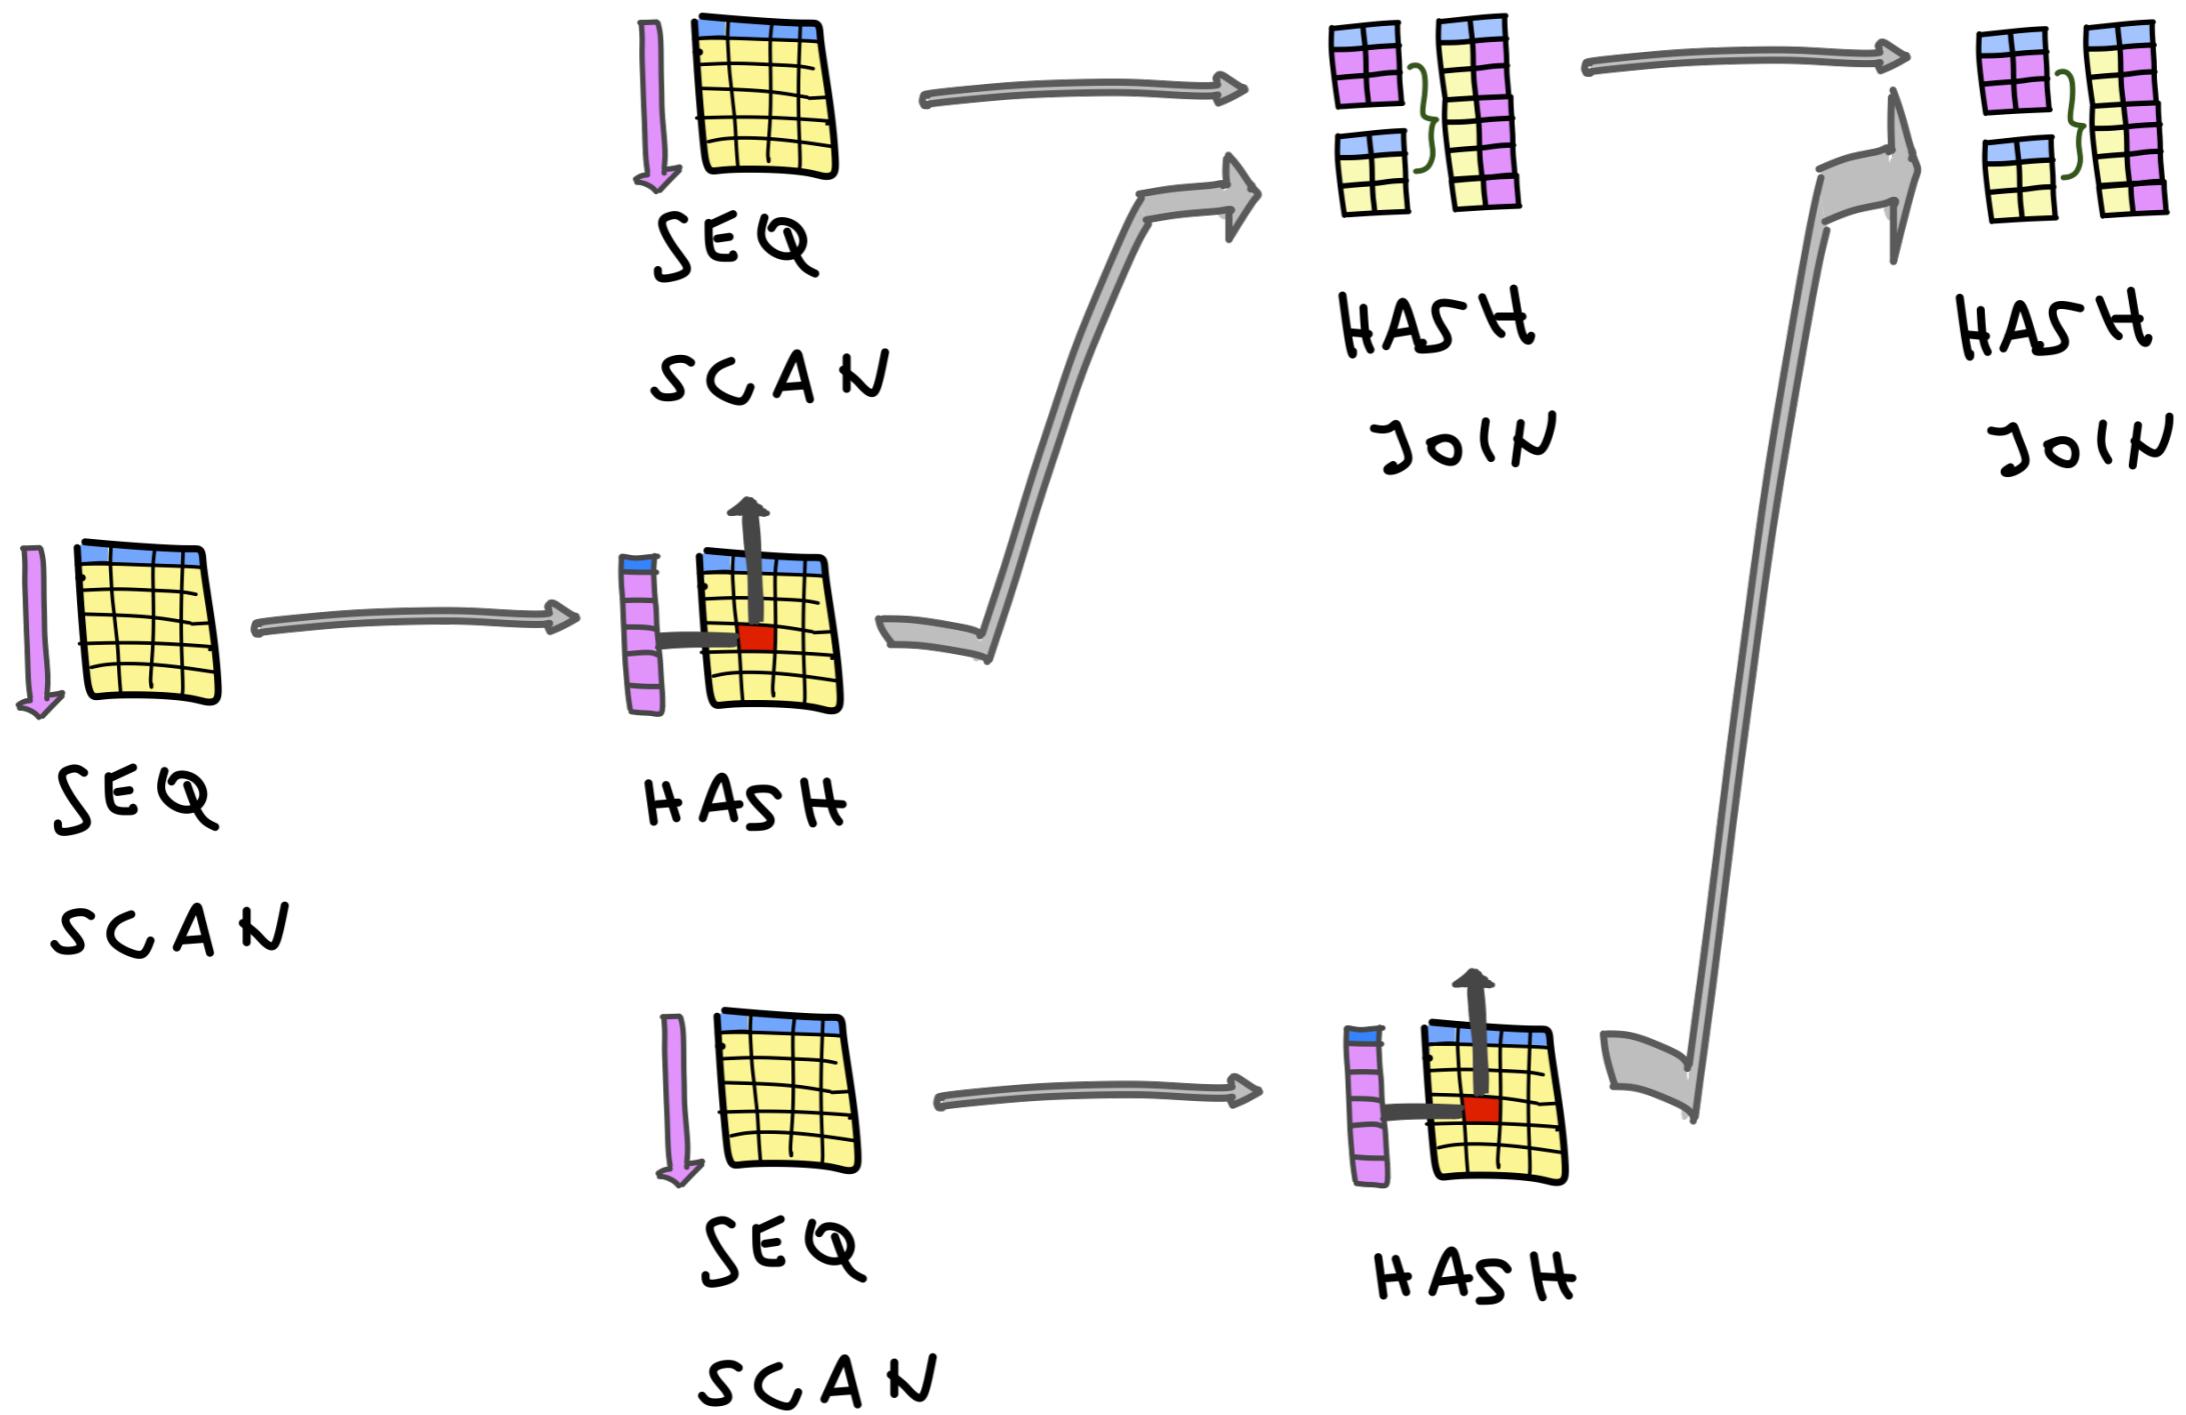
```
      -> Index Scan using index_paragraph_on_charid
on paragraph p  (cost=0.29..7913.89 rows=35465
width=361)
```

# MULTI JOIN

---

```
select *
  from character c
join character_work cw on c.charid = cw.charid
join work w on w.workid = cw.workid;
```

```
Hash Join (cost=43.45..102.93 rows=1346 width=173)
  Hash Cond: ((cw.workid)::text = (w.workid)::text)
-> Hash Join (cost=41.48..82.45 rows=1346 width=74)
    Hash Cond: ((cw.charid)::text =
(c.charid)::text)
      -> Seq Scan on character_work cw
(cost=0.00..22.46 rows=1346 width=19)
      -> Hash (cost=25.66..25.66 rows=1266 width=55)
          -> Seq Scan on "character" c
(cost=0.00..25.66 rows=1266 width=55)
      -> Hash (cost=1.43..1.43 rows=43 width=99)
          -> Seq Scan on work w (cost=0.00..1.43 rows=43
width=99)
```



# QUERY PLANS EXPLOSION

---

- Each join:
  - 3 join strategies
  - $x^2$  driving/inner table combinations
  - seq scan vs. index access (3 types)
  - table permutations
- ...
- Explosion of possible query plans

# GENETIC OPTIMIZER

---

1. Generate a set of initial population - query plans
  2. Cross and mutate
  3. Evaluate fitness (plan cost) and select the most fit plans (lowest cost)
  4. Continue at 2. until stop condition
- 
- Activated based on the number of joined tables (`show geqo_threshold;`)
  - For further reading <https://www.postgresql.org/docs/current/geqo.html>

# OTHER TYPES OF JOINS

---

```
select * from paragraph p
  left join character c on p.charid = c.charid
```

**Hash Left Join** (cost=41.48..2623.78 rows=35465 width=416)

  Hash Cond: ((p.charid)::text =  
  (c.charid)::text)

    -> Seq Scan on paragraph p  
(cost=0.00..2094.65 rows=35465 width=361)

    -> Hash (cost=25.66..25.66 rows=1266 width=55)

        -> Seq Scan on "character" c  
(cost=0.00..25.66 rows=1266 width=55)

# OTHER TYPES OF JOINS

---

- Hash Left Join, Hash Right Join, Merge Left Join, Merge Right Join, Nested Loop Left Join, Nested Loops Right Join...
- Semijoins, Antijoins, ...

# AGGREGATIONS

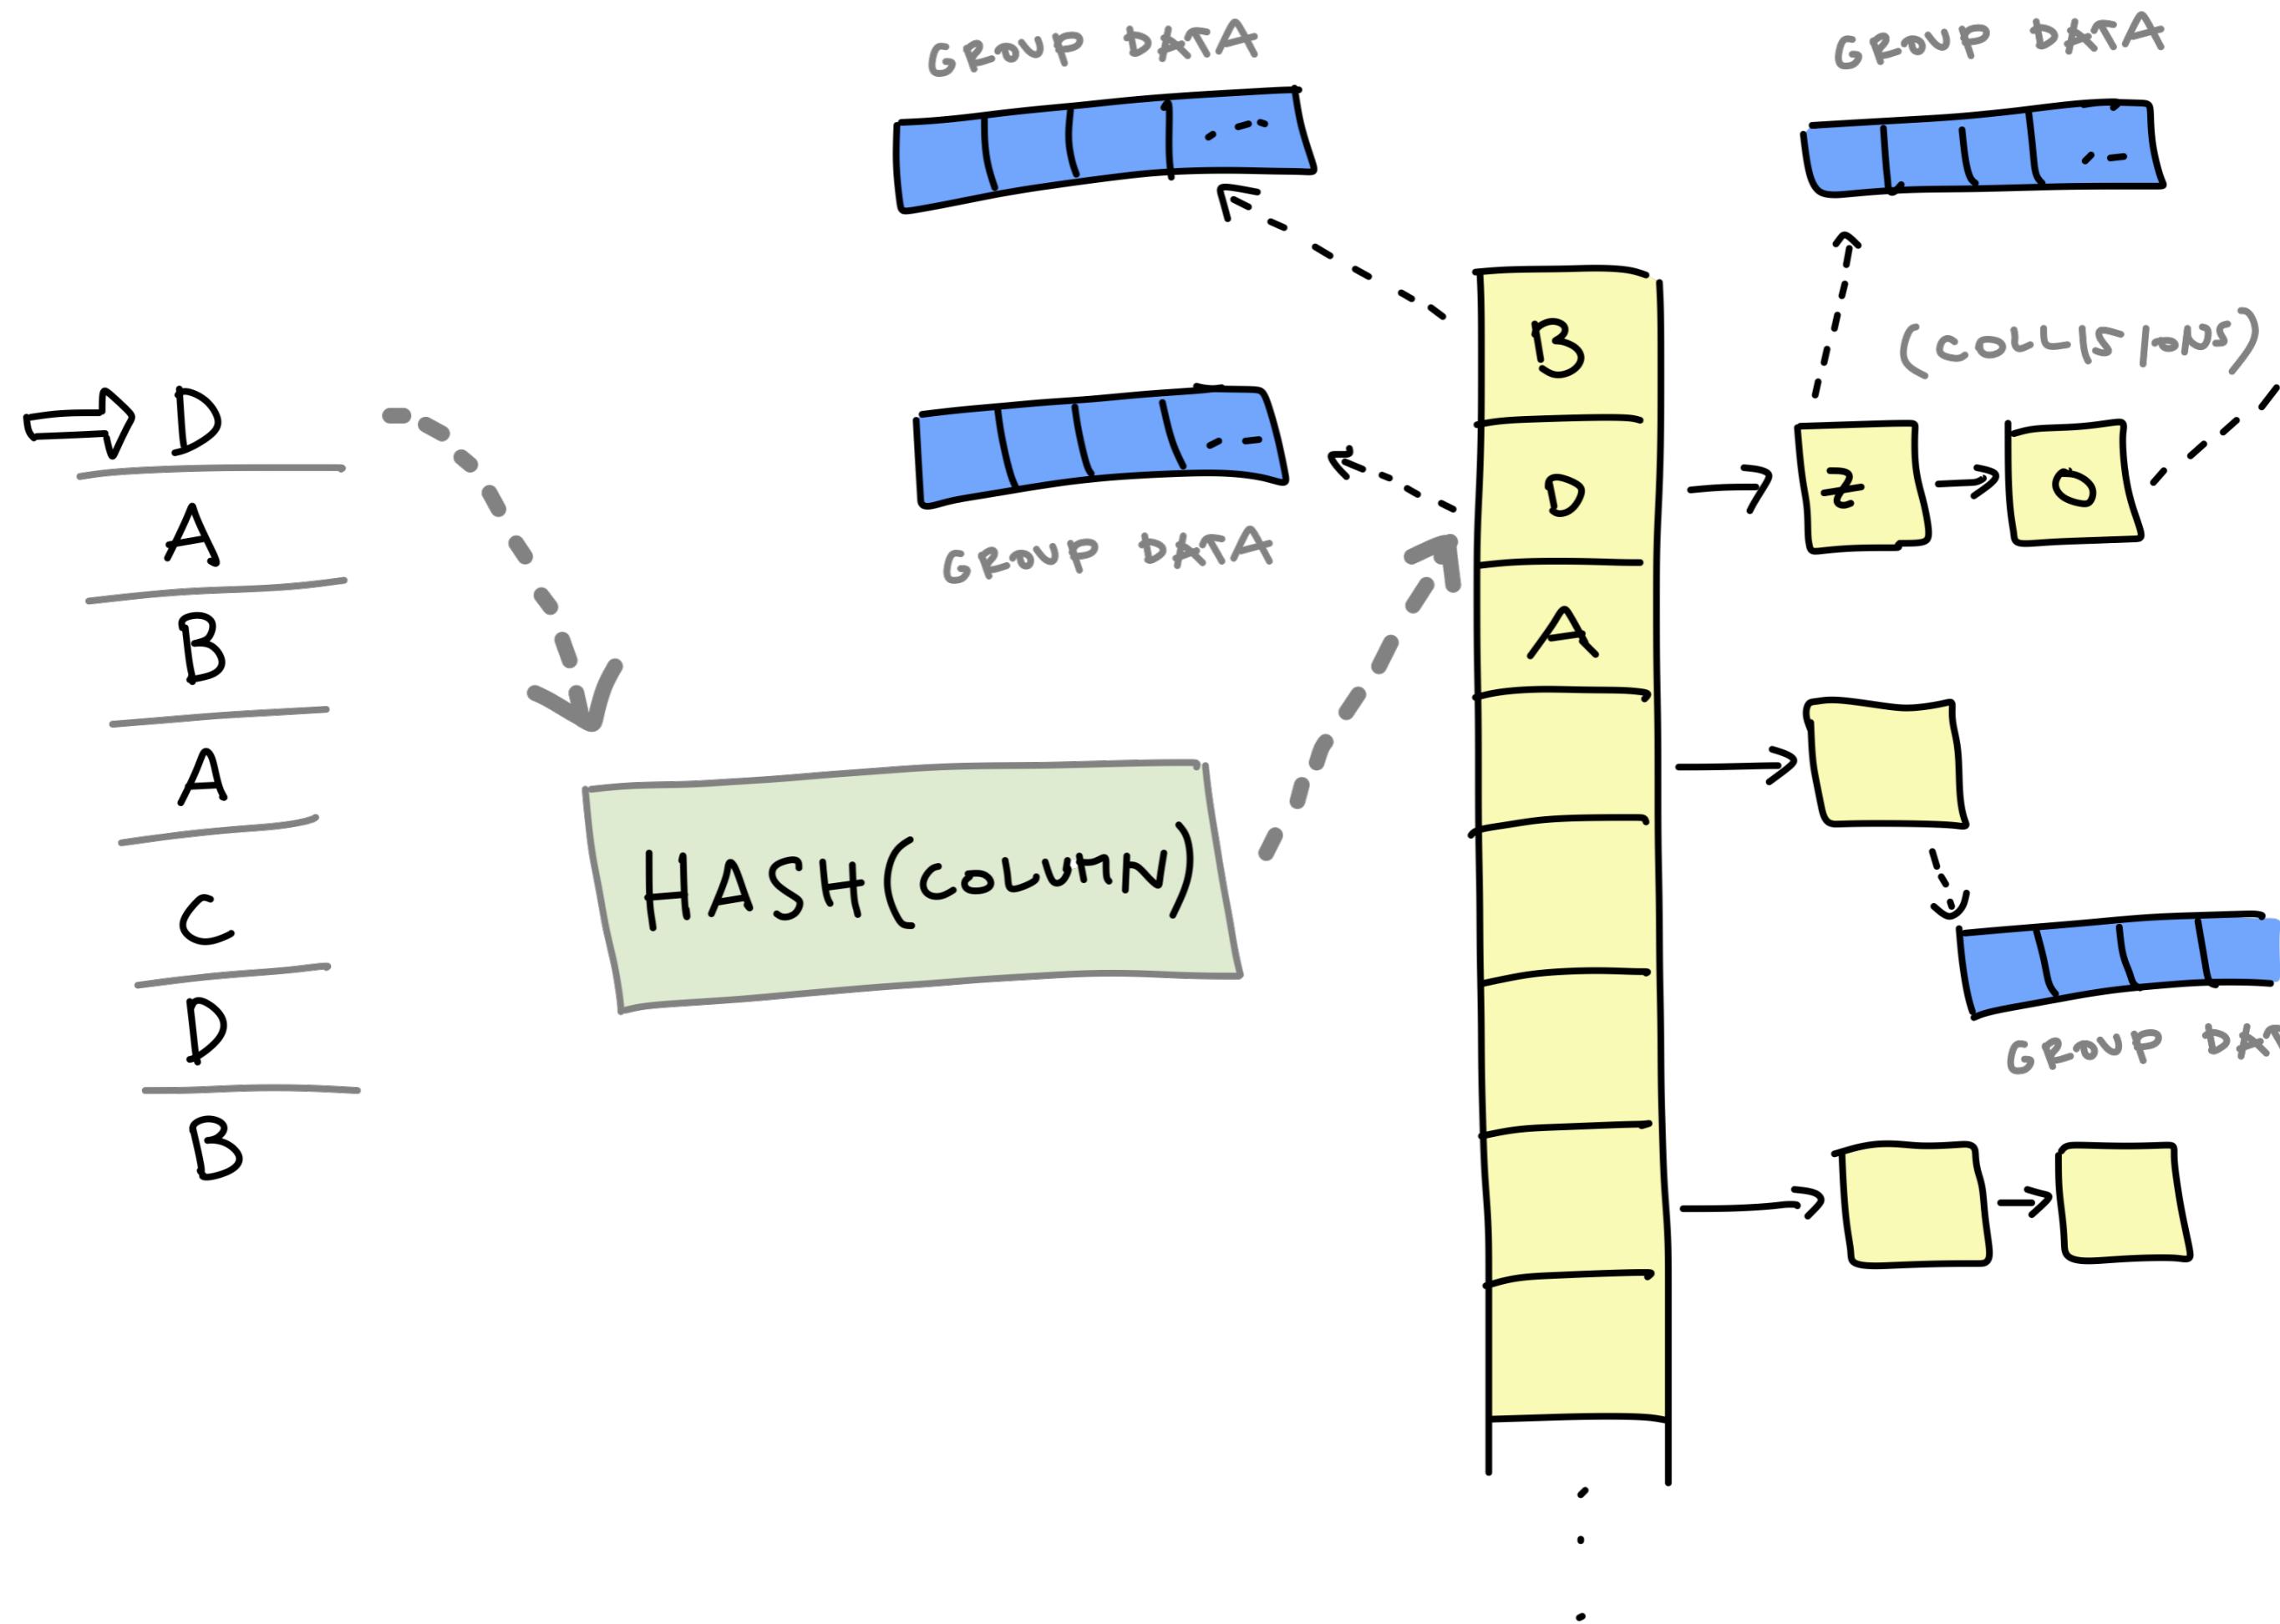
# HASHAGGREGATE

---

```
explain select department, count(*) from  
documents group by department;
```

```
HashAggregate (cost=18276.36..18277.14 rows=78  
width=42)
```

```
    -> Seq Scan on documents  
(cost=0.00..16520.24 rows=351224 width=42)
```



# GROUP/GROUPAGGREGATE

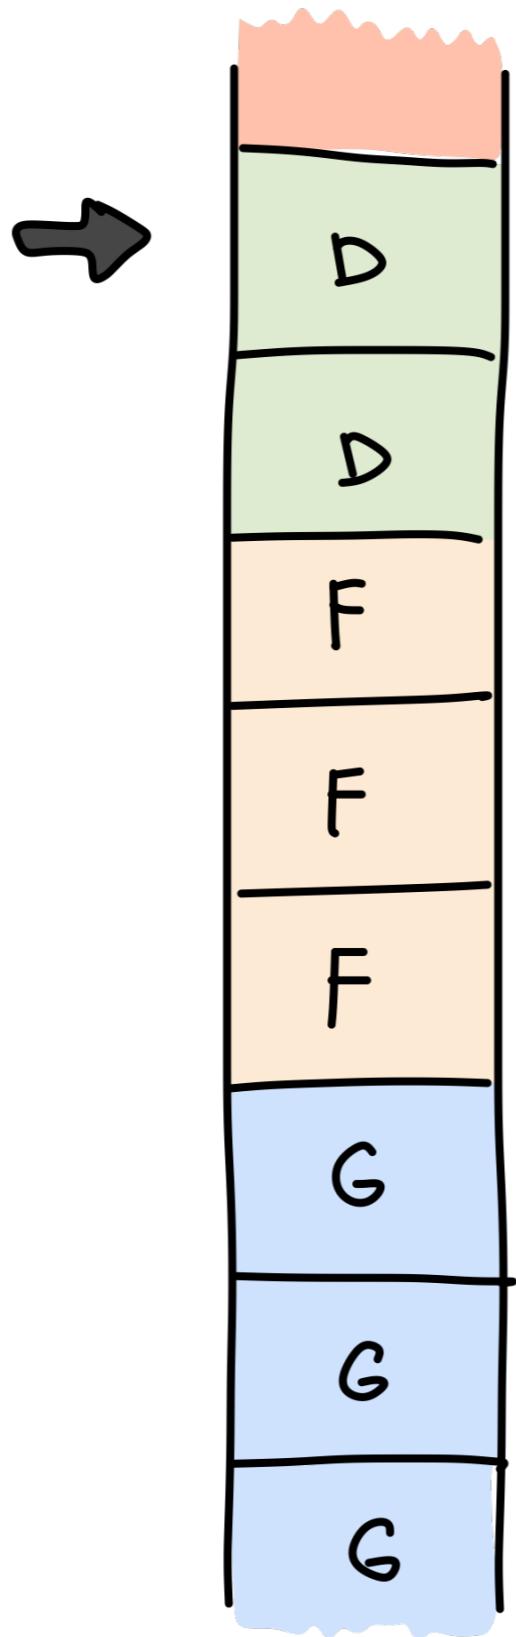
---

```
create index index_documents_on_dc on
documents(department, customer); vacuum analyze;
explain select customer, count(*) from documents
where department = 'Rozhlas a televizia
Slovenska' group by customer;
```

**GroupAggregate** (cost=0.55..1287.26 rows=335 width=41)

  -> **Index Only Scan** using  
index\_documents\_on\_dc on documents  
(cost=0.55..1200.50 rows=16683 width=41)

    Index Cond: (department = 'Rozhlas a  
televizia Slovenska'::text)



# SUMMARY (WHAT YOU SHOULD KNOW)

---

- Understand the difference between 1 composite index and several single-column indices, draw a query plan for simple query with bitmap index scans (bitmap and/or).
- Understand why filtering on non-leading column in composite index is not efficient. Understand why like '%something%' is inefficient.
- Understand 3 join types and how they work. Estimate for a given query and schema, which join type would be most efficient.