



ADVANCED DATATYPES

Array, JSON and GIN indices

ARRAYS

ARRAYS

```
create table flights(  
    airport_from text,  
    airport_to    text,  
    services      text[]  
) ;
```

```
insert into flights values('BTS', 'VIE',  
'{lunch, 1st class}');  
insert into flights values('BTS', 'FRA',  
'{lunch, wifi}');
```

ARRAYS

```
select *  
from flights f  
where 'wifi' = any(f.services);
```

```
select *  
from flights f  
where f.services @> '{wifi}';
```

USAGE?

- Denormalized schema (not necessarily bad)
- Convenience (avoids many-to-many)
- Keeping track of state in recursive queries (to avoid cycles)

INDEXING ARRAY ACCESS

```
create index idx_services on
flights(services); -- b-tree
```

```
set enable_seqscan = 'off';
-- we have only 2 rows
```

INDEXING ARRAY ACCESS

```
explain  
select *  
from flights f  
where 'wifi' = any(f.services);
```

```
Seq Scan on flights f  (cost=0.00..1.07  
rows=1 width=96)  
  Filter: ('wifi'::text = ANY (services))
```

INDEXING ARRAY ACCESS

```
explain  
select *  
from flights f  
where f.services = '{lunch, wifi}';
```

Index Scan using idx_services on flights f
(cost=0.13..8.15 rows=1 width=96)
 Index Cond: (services =
'{lunch,wifi}':text[])

INDEXING ARRAY ACCESS

```
drop index idx_services;
create index idx_services on flights using
gin(services); -- generalized inverted index

set enable_seqscan = 'off';
-- we still have only 2 rows
```

INDEXING ARRAY ACCESS

```
explain  
select *  
from flights f  
where f.services @> '{wifi}';
```

```
Bitmap Heap Scan on flights f  
(cost=8.00..12.01 rows=1 width=96)  
  Recheck Cond: (services @>  
  '{wifi}'::text[] )  
    -> Bitmap Index Scan on idx_services  
(cost=0.00..8.00 rows=1 width=0)  
      Index Cond: (services @>  
  '{wifi}'::text[] )
```

INVERTED INDEX

TABLE

TID	SERVICES
1	['LUNCH', '1ST CLASS']
2	['LUNCH', 'WIFI']

INVERTED INDEX

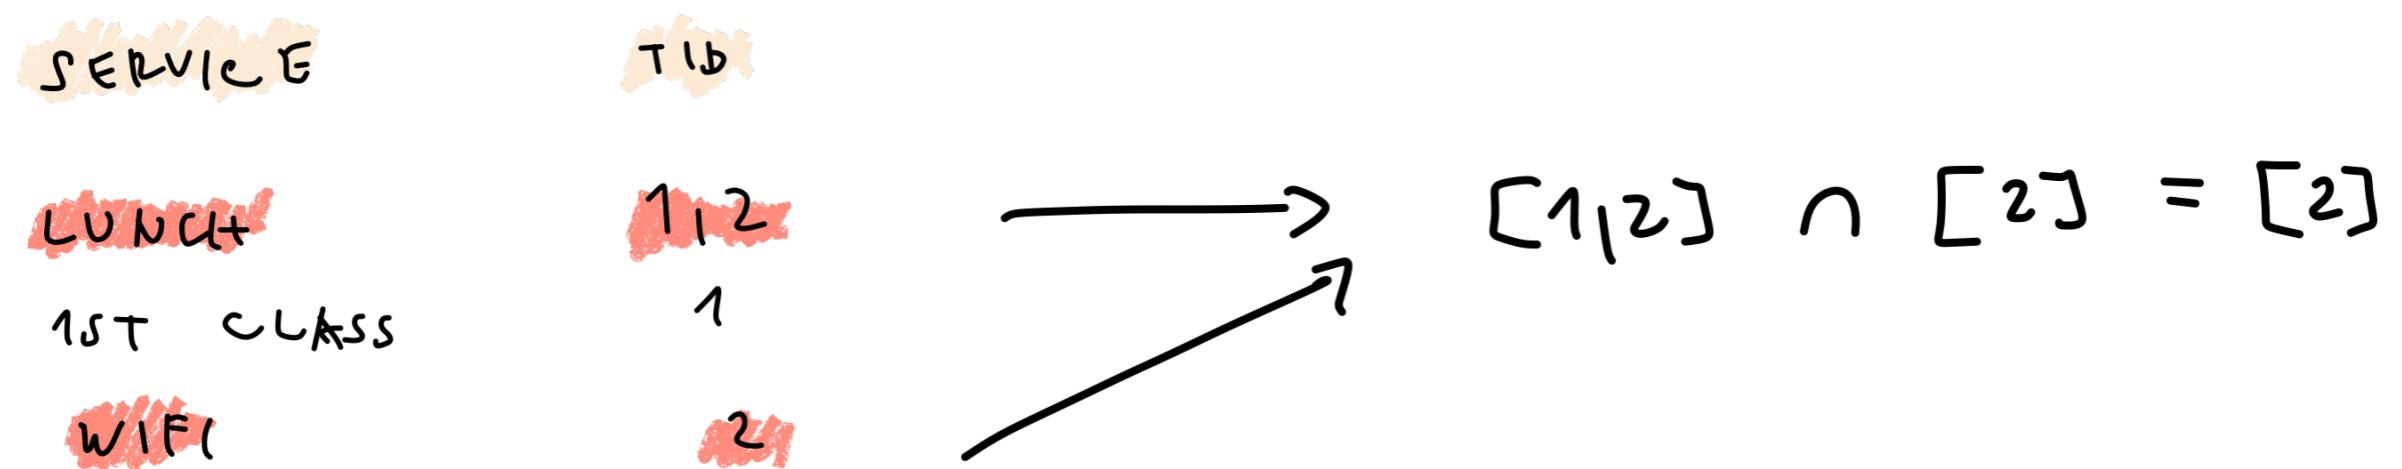
SERVICE	TID
LUNCH	1, 2
1ST CLASS	1
WIFI	2

INDEXING ARRAY ACCESS

```
explain  
select *  
from flights f  
where f.services @> '{wifi, lunch}';
```

INVERTED INDEX

INVERTED INDEX



TRANSFORM ARRAY BACK TO ROWS (E.G. TO JOIN)

```
select airport_from, airport_to,  
unnest(services)  
from flights f  
where 'wifi' = any(f.services);
```

airport_from	airport_to	unnest
BTS	VIE	lunch
BTS	VIE	wifi

JSON

JSON

- JavaScript Object Notation
- Very popular
 - Simple, convenient, human readable, universal support
- Basic data types

JSON EXAMPLE

```
{  
  "name": "Advanced datatypes",  
  "length": 120,  
  "topics": ["arrays", "gin", "json"],  
  "teacher": {  
    "name": "Tomas",  
    "surname": "Kramar"  
  }  
}
```

USAGE?

- Data schema not known in advance
- Impractical to keep adding columns
- Can introduce new keys (even nested) anytime
- Checks JSON validity

```
.....  
create table lectures_json(doc json);  
  
insert into lectures_json values('{name:  
"Advanced datatypes", "length": 120}');  
  
ERROR: invalid input syntax for type json  
LINE 1: insert into lectures_json  
values('{name: "Advanced datatypes...  
DETAIL: Token "name" is invalid.  
CONTEXT: JSON data, line 1: {name...  
.....
```

```
.....  
select doc->name from lectures_jsonb;  
ERROR: column "name" does not exist
```

```
select doc->'name' from lectures_json;
```

?column?

"Advanced datatypes"

2 JSON DATATYPES IN POSTGRES

- JSON
 - text representation
 - preserves formatting and key ordering
 - must be re-parsed on each read
- JSONB
 - binary representation
 - loses formatting and key ordering
 - slower to write (encoding), faster to read (no re-parsing)
 - supports indexing

```
.....  
create table lectures_json(doc json);  
create table lectures_jsonb(doc jsonb);  
  
insert into lectures_json values('{"name":  
"Advanced datatypes", "length": 120}');  
  
insert into lectures_jsonb  
values('{"name": "Advanced datatypes",  
"length": 120}');
```

```
.....  
select * from lectures_json;  
{"name": "Advanced datatypes", "length":  
120}
```

```
select * from lectures_jsonb;  
{"name": "Advanced datatypes", "length":  
120}
```

INDEXING JSONB

```
create index idx_doc on lectures_jsonb using  
gin(doc);
```

-- Allows "containment" queries

```
explain select * from lectures_jsonb  
where doc @> '{"name": "Advanced datatypes"}';
```

Bitmap Heap Scan on lectures_jsonb
(cost=12.00..16.01 rows=1 width=32)

 Recheck Cond: (doc @> '{"name": "Advanced
datatypes"}'::jsonb)

 -> **Bitmap Index Scan** on idx_doc
(cost=0.00..12.00 rows=1 width=0)

 Index Cond: (doc @> '{"name": "Advanced
datatypes"}'::jsonb)

GIN INDEX OVER JSONB

TID

1

JSONB

```
{ "NAME": "ADVANCED DATA ..", "LENGTH": 120 }
```

2

```
{ "NAME": "GIS", "LENGTH": 120 }
```

KEY / VALUE

NAME (ADVANCED DATA ..)

LENGTH (120)

NAME (GIS)

TID

1

112

2

HANDLING NESTING BY FLATTENING

```
{
```

```
"TEACHER": {
```

```
  "NAME": "Thomas"
```

```
}
```

```
}
```

```
TEACHER.NAME / Thomas
```

INDEXING JSONB OVER A SINGLE "PROPERTY"

```
create index idx_lectures_on_name on  
lectures_jsonb using gin((doc->'name'));
```

```
explain select * from lectures_jsonb  
where doc->'name' ? 'Advanced datatypes';
```

Bitmap Heap Scan on lectures_jsonb
(cost=8.00..12.02 rows=1 width=32)
 Recheck Cond: ((doc -> 'name'::text) ?
 'Advanced datatypes'::text)
 -> **Bitmap Index Scan** on idx_lectures_on_name
(cost=0.00..8.00 rows=1 width=0)
 Index Cond: ((doc -> 'name'::text) ?
 'Advanced datatypes'::text)

JSON FUNCTIONS

- `row_to_json`
 - `select row_to_json(t) from (select name, length from lectures) t`
- `array_agg`
- `array_to_json`
- `::json cast`

WHAT YOU SHOULD KNOW

- Understand principles of inverted index and how it's different from b-tree
- Understand how inverted index answers compound conditions (@> {a, b})
- Understand how b-tree index works for arrays and its limitations compared to inverted index