



RECURSIVE SQL

2 DEMOS

- Factorial "table" - not a real-world use-case, but easy
- Optimal mass transport transfers (simple version of imhd.sk)
- real-world, but hard(er)

**BREAKING NEWS:
THERE IS NO RECURSION
IN SQL**

RECURSIVE SQL = ITERATIVE SQL

- Traditional recursion:
 - Start with unknown, back down until known value
- Recursive Iterative SQL
 - Start with known, build up the answer
 - Get ready for weird syntax

DEMO: FACTORIAL

```
with recursive fact(n, fact) as (
    select 0, 1
    union all
    select f.n+1, f.fact * (f.n+1) from fact f
)
select * from fact limit 10
```

#1

```
with recursive fact(n, fact) as (
    select 0, 1
    union all
    select f.n+1, f.fact * (f.n+1) from fact f
)
select * from fact limit 10
```

WORKING TABLE

WORKING TABLE / NEXT

OUTPUT TABLE

0 1

0 1

2

```
with recursive fact(n, fact) as (
    select 0, 1
    union all
    select f.n+1, f.fact * (f.n+1) from fact f
)
select * from fact limit 10
```

WORKING TABLE

0 1

WORKING TABLE / NEXT

1 1

OUTPUT TABLE

0 1
1 1

3

with recursive fact(n, fact) as (
 select 0, 1
 union all
 select f.n+1, f.fact * (f.n+1) from fact f
)
select * from fact limit 10

WORKING TABLE

1 1

WORKING TABLE / NEXT

2 2

OUTPUT TABLE

0 1
1 1
2 2

4

```
with recursive fact(n, fact) as (
    select 0, 1
    union all
    select f.n+1, f.fact * (f.n+1) from fact f
)
select * from fact limit 10
```

WORKING TABLE

2 2

WORKING TABLE / NEXT

3 6

OUTPUT TABLE

0 1

1 1

2 2

3 6

UNION / UNION ALL

- Union = Append to output table
- Union all = Add to output table but avoid duplicate rows

DEMO: IMHD

```
with recursive hops(start_name, line, end_name, end_stop_id,
stops_visited, lines_used, duration) as (
    select s.name, c.line, e.name, e.id, array[s.name, e.name],
array[c.line], c.duration
        from connections c
        join stops s on c.start_stop_id = s.id
        join stops e on c.end_stop_id = e.id
    where s.name = 'Nad lúčkami'
        and c.line not like 'N%'
union
    select h.end_name, c.line, e.name, e.id, h.stops_visited || e.name,
           case when c.line = h.line then h.lines_used else h.lines_used ||
c.line end, h.duration + c.duration
        from hops h
        join connections c on h.end_stop_id = c.start_stop_id
        join stops e on c.end_stop_id = e.id
    where c.line not like 'N%'
        and (c.line = h.line or not (c.line = any(h.lines_used)))
        and not (e.name = any(h.stops_visited))
        and h.end_name != 'Zochova'
        and array_length(h.lines_used, 1) <= 2
)
select h.stops_visited, h.lines_used, h.duration from hops h
where h.end_name = 'Zochova'
order by array_length(h.lines_used, 1) asc
```

BUT WHY?

- Recursive query
 - algorithm + data are collocated = fast processing
- Recursive algorithm pulling data from the database
 - need to pull data from db to app = slow(er) processing
- Graph database (e.g. Neo4j)
 - may be an overkill
 - need to synchronize data

WHAT YOU SHOULD NOW

- Write a simple recursive query ("1-2 cups of coffee"-level in labs)
- Understand when to use recursive SQL and when to use a different approach